SERA

Personal AI Career Agent

Built on Claude Code · Gmail MCP · SQLite · Local Dashboard

What Is SERA?

A locally-run AI agent for end-to-end job search workflow automation

AI-Powered

Built on Claude Code with specialized markdown instruction files (agent files) per workflow role

Persistent Memory

SQLite databases track every reviewed role, application status, scores, and session telemetry across time

Hard Rules + Soft Judgment

Deterministic auto-filters (Verizon restriction, comp floor) combined with evaluative AI scoring

Live Dashboard

Python-served local web UI at localhost:5500 visualizing the full pipeline in real time

5
Agent Files

3
Data Stores

1
Live Dashboard

∞
Roles Tracked

System Components

Five agent files — each with a single, well-defined responsibility

`sera.md` Routing Layer	Entry point — reads intent, routes to agent, runs DB health check, writes session logs
`email-scanner.md` Email Scanner	Gmail MCP → classify → extract jobs → score vs. two resumes → tiered output → DB write
`job-match.md` Job Match Reviewer	Full JD parse → score → TL;DR verdict → optional tailored resume → DB + flat-file backup
`job-log.md` Job Logger	NJ UI compliance CSV — triggered on application confirm, no prompt (legally required)
`config.md` Config / SSOT	File paths, DB schema, score thresholds, auto-filter rules, comp floor, commute radius

Data Layer

Persistent storage across five purpose-built stores

job-tracker.db

SQLite

Primary — all reviewed postings, scores, status, comp ranges, links. The system of record.

monitor.db

SQLite

Agent monitoring — session log, event log, scan metrics, quality flags, error telemetry.

reviewed-postings.md

Markdown

Flat-file human-readable backup of every review. Portable, no DB required to read.

job-search-log.csv

CSV

NJ UI compliance — legally required work search record, auto-appended on each application.

tailored/ directory

Files

Per-role tailored resumes generated on request during job-match review. Named by role + date.

What Makes It Agentic

Six capabilities that separate a workflow tool from an AI agent

Persistent Memory

Knows what's been reviewed, applied to, and rejected — across sessions, without re-prompting

Duplicate Detection

Checks DB before every review — flags previously seen postings before wasting a review cycle

Multi-Step Reasoning

Parse → assess → score → recommend → optionally revise resume — chained, not isolated

Tool Use

Gmail MCP for retrieval, Bash for DB reads/writes, file I/O for resume and log management

Hard Rules + Soft Judgment

Auto-filters are deterministic; scoring is evaluative — the system knows the difference

Self-Monitoring

Writes session telemetry to monitor.db, surfaces quality flags when something goes wrong

Routing Flow

sera.md dispatches every session — three entry paths, two shared outputs

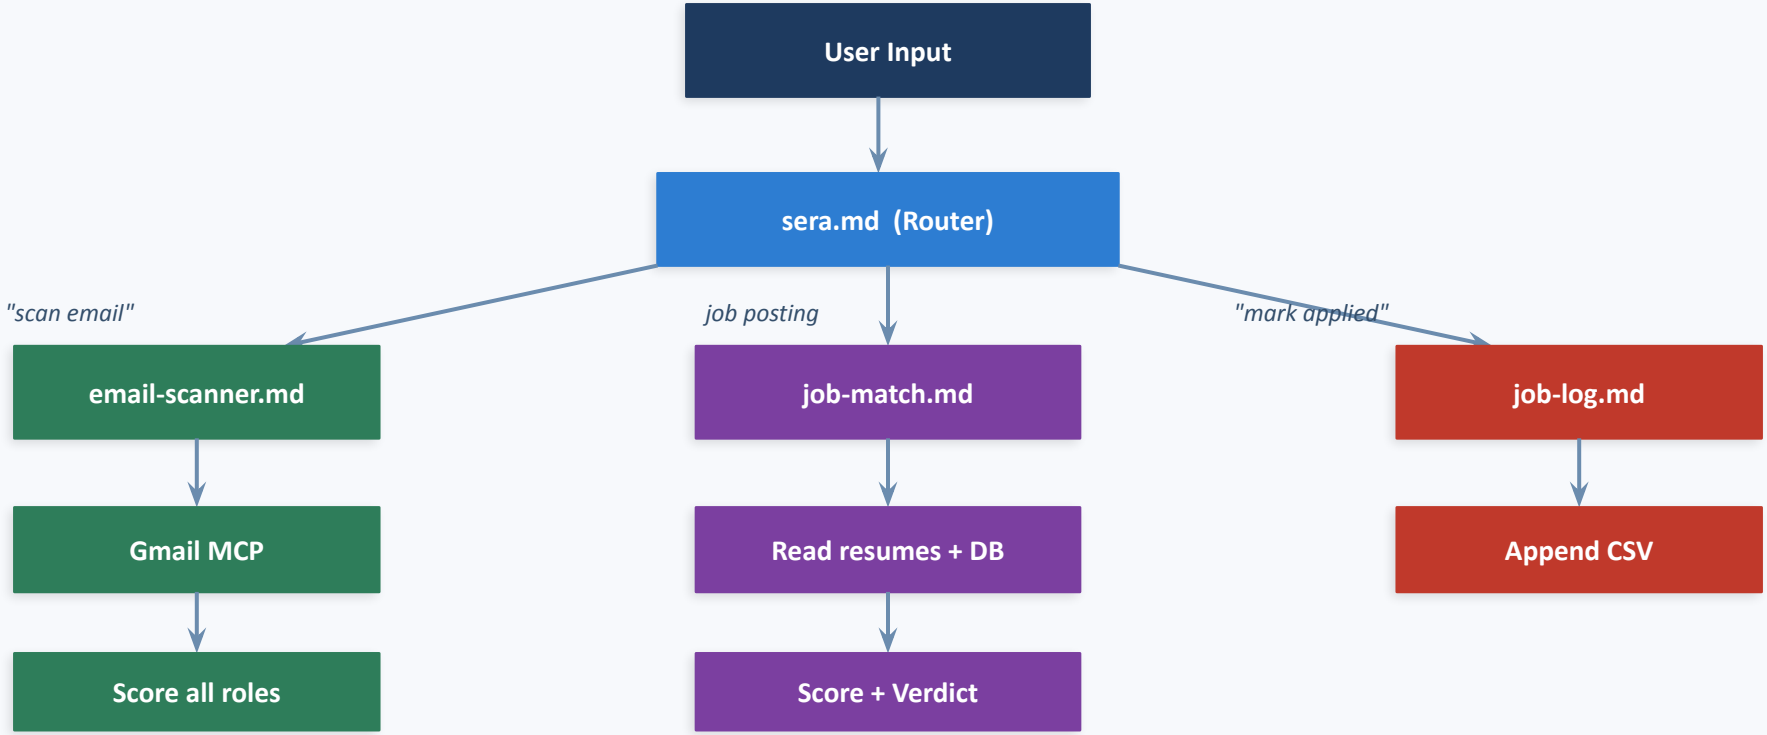

All writes → job-tracker.db + reviewed-postings.md

Dashboard → localhost:5500

Built to Learn the System

SERA reflects a career-long pattern: understanding how systems actually work from the inside, then building logic to make them run better. The job search is a pipeline. This is the automation layer.

- Locally run. No third-party data exposure.
- Modular. Each agent file is replaceable or extensible.
- Legally compliant. NJ UI work search record built in.